Rethinking Scientific Analysis Using Modern Scalable Systems

ABSTRACT

Revolutions in data acquisition are drastically changing how science conducts experiments. For example, "next-generation" sequencing technologies have driven exponential growth in the volume of genomic data, and similar trends impact many fields which rely on imaging, such as astronomy and neuroscience. Although there have been early attempts to use MapReduce systems to accelerate the processing of these datasets, they have conceded efficiency in favor of using legacy storage formats and software.

Since the amount of scientific data that is being captured is increasing exponentially, we have a good opportunity to rethink how we process and store these datasets. In this paper, we introduce a set of principles for implementing scientific analyses efficiently using commodity "big data" systems. We motivate these principles with an example genomics pipeline which leverages open-source MapReduce and columnar storage techniques to achieve a $> 50\times$ speedup over traditional genomics systems, at half the cost. To demonstrate the generality of our architecture, we then achieve a $2.5\times$ improvement over the state-of-the-art MPI-based system for an astronomy task.

Categories and Subject Descriptors

L.4.1 [**Applied Computing**]: Life and medical sciences—*Computational biology*; H.1.3.2 [**Information Systems**]: Data management systems—*Database management system engines, parallel and distributed DBMSs*; E.3.2 [**Software and its Engineering**]: Software creation and management—*Software Development Process Management*

General Terms

Design

Keywords

Analytics, MapReduce, Genomics, Scientific Computing

1. INTRODUCTION

Major improvements in scientific data acquisition techniques have increased scientific data storage and processing needs [34, 8]. In fields like neuroscience [14] and genomics [39], particle physics, and astronomy scientists routinely perform experiments that use terabytes (TB) to petabytes (PB) of data. While traditional scientific computing platforms are optimized for fast linear algebra, many emerging domains make heavy use of statistical learning techniques and user defined functions (UDFs) on top of semistructured data. This move towards statistical techniques has been driven by the increase in the amount of data available to scientists, as well as the rise of statistical systems which are accessible to non-experts, such as `Scikit-learn` [32] and `MLI` [38].

While the increase in the amount of scientific data available is a boon for scientists, it puts significant stress on existing software. Using the current "best practice" genomic variant calling pipeline [4], it takes approximately 120 hours to process a single, high-quality human genome using a single, beefy node [40]. To address these challenges, scientists have started to apply computer systems techniques such as MapReduce [27, 35, 22] and columnar storage [15] to custom scientific compute/storage systems. While these systems have improved the analysis cost and performance, naïve implementations incur significant overheads due to constraints of the legacy formats and codebases that they use.

Since the amount of scientific data being generated is growing so quickly, we cannot afford to be saddled by legacy software and formats. New scientific projects such as the 100K for UK, which aims to sequence the genomes of 100,000 individuals in the United Kingdom [13] will generate three to four *orders of magnitude* more data than prior projects like the 1000 Genomes Project [37]. While it is important to still be able to use data stored in legacy formats, the massive amount of *new* data provides us with an opportunity to rethink how we compose our systems. By choosing the correct commodity computing systems, we can provide better performance and scalability than custom systems while enhancing the abstractions exposed to scientists.

In this paper, we demonstrate a system built using Apache Avro, Parquet, and Spark [2, 3, 46], that achieves a $50\times$ increase in throughput over the current best practice pipeline for processing genomic data, while reducing the analysis cost by 50%. In the process of creating this system, we developed

a "narrow waisted" layering model for building similar scientific analysis systems. This narrow waisted model is inspired by the OSI model for networked systems [47]. We demonstrate the generality of this model by using it to implement a system for processing astronomy images, and achieve a $> 3\times$ performance improvement over a state-of-the-art Message Passing Interface (MPI) based pipeline.

An additional benefit of our stack model is that it addresses a subtle problem in earlier scientific processing and storage systems. In these systems, characteristics of the data format on disk would bleed into the computing model. For example, in the current Sequence/Binary Alignment and Map (SAM/BAM [24]) formats for storing genomic alignments, constraints about record ordering are required in order to enable specific computing patterns. We believe that this is an abstraction inversion, which makes it difficult to perform some other access patterns. In §3, we elucidate why this is a significant problem, and then in §5, we then introduce efficient algorithms for supporting these computational patterns without forcing constraints on the storage layer.

To summarize, our paper makes the following contributions:

1. We characterize the important data and application patterns of emerging scientific domains, and review why current scientific software does not fully satisfy these use cases.

2. We develop a modular, technology independent stack model to guide the design and implementation of scientific processing and storage systems. We then demonstrate how this stack model resolves both design anti-patterns *and* functional bugs in current systems.

3. We implement a genomics pipeline which achieves a $50\times$ performance improvement and a $2\times$ cost improvement over current "best practice" systems, and an astronomy image processing pipeline which provides a $3\times$ improvement over the current "state-of-the-art" MPI system.

4. Finally, we illustrate how to provide coordinate-system guided processing in parallel across shared-nothing systems. These traversals are necessary in many scientific applications.

All of the software (source code and executables) described in this paper is made available free of charge under the permissive Apache 2 license.

2. BACKGROUND

As our work exists at the intersection of computational science and data management and processing systems, our architectural approach is informed by recent trends in both areas. The design of large scale data management has changed dramatically since the landmark papers by Dean and Ghemawat [9, 10] which described Google's `MapReduce` system. Over a similar timeframe, scientific fields have moved to take advantage of improvements in data acquisition technologies. For example, since the Human Genome Project finished in 2001 [21], the price of genomic sequencing has dropped by $10,000\times$ [29]. This drop in cost has enabled the capture of petabytes of sequence data, which has (in turn) enabled significant population genomics experiments like the 1000 Genomes project [37], and The Cancer Genome Atlas (TCGA, [42]). These changes are not unique to genomics; indeed, fields such as neuroscience [8] and astronomy [41] are experiencing similar changes.

Although there has been significant progress in the development of systems for processing large datasets (e.g., the development of first generation MapReduce systems [9], followed by iterative MapReduce systems like Spark [46], as well as parallel and columnar DBMS [1, 20]), the uptake of these systems in the scientific world has been slow. Most implementations have either used MapReduce as an inspiration for programming API design [27], or have been limited systems which have used MapReduce to naïvely parallelize existing toolkits [22, 35]. These approaches are perilous for several reasons:

- A strong criticism levied against the map-reduce model is that the API is insufficiently expressive for describing complex tasks. As a consequence of this, tools like the GATK [27] which adopt MapReduce as a programming model (but *not an execution strategy*) force significant restrictions on algorithm implementors. For example, a GATK `walker` is provided with a single view over the data (a sorted iterator over a specified region), and is allowed limited reduce functionality.

- A major contribution of systems like MapReduce [10] and Spark [46, 45] is the ability to reliably distribute parallel tasks across a cluster in an automated fashion. In practice, to run tools like the GATK across a cluster, organizations have rolled their own systems for sharding and persisting intermediate data, and managing failures and retries. This is not only an inefficient duplication of work, but it is also a source of inefficiency during execution: the performance of iterative stages in the GATK is bottlenecked by I/O performance. Additionally, the sharding techniques used are based on prior knowledge of the structure of the genome and limit scale-up to a $10\times$ speedup on 23 machines.

- The naïve Hadoop-based implementations in Crossbow [22] and Cloudburst [35] lead to good speedups, but add significant overhead. Several of the methods that they parallelize incur high overhead due to duplicated loading of indices (for fast aligners, loading of large indices can be a primary I/O bottleneck) and poor broadcasting of data.

A notable exception to the above criticisms is the `Thunder` system, which was developed for processing terabytes of neuroscience imaging data [14]. `Thunder` performs a largely statistical workload, and the significant tasks in terms of execution time are clustering and regression. The system is constructed using Spark and Python and is designed to process datasets larger than 4 TB, and leverages significant functionality from the MLI/MLLib libraries [38]. Unlike pre-existing pipelines, they are able to use Spark's in-memory caching to amortize load time across several pipeline stages.

Additionally, they use Spark's filtering primitives to allow scientists to cut problems into subproblems. This is a common trend across scientific analyses, and is one of the reasons that we advocate for the use of a columnar store with efficient predicate pushdown.

Recent work by Diao et al [12] has looked at optimizations to MapReduce systems for processing genomic data. They borrow strategies from the query optimization literature to reorder computation to minimize data shuffling. While this approach does improve shuffle traffic, several preprocessing stages cannot be transposed. For instance, reversing the order of indel realignment and base quality score recalibration (see §4.1) will change the inferred quality score distribution. Additionally, we believe that the shuffle traffic that Diao et al observe is an artifact caused by the abstraction inversion discussed in §1. As we demonstrate in §4.1, these penalties can be eliminated by restructuring the preprocessing algorithms.

One interesting trend of note is the development of databases specifically for scientific applications. The exemplar is SciDB, which provides an array based storage model as well as efficient linear algebra routines [6]. While arrays accelerate many linear algebra based routines, they are not a universally great fit. For many genomics workloads, data is semistructured and may consist of strings, boolean fields, and an array of tagged annotations. Other systems like the Genome Query Language [19] have extended SQL to provide efficient query semantics across genomic coordinates. While GQL achieves performance improvements of up to $10\times$ for certain algorithms, SQL is not an attractive language for many scientific domains, which make heavy use of user designed functions, which may be cumbersome in SQL.

One notable area where database techniques have been leveraged by scientists is in the data storage layer. Due to the storage costs of large genomic datasets, scientists have introduced the CRAM format which uses columnar storage techniques and special compression algorithms to achieve a 30% reduction in size over the original BAM format [15]. While CRAM achieves good compression, it imposes restriction on the ordering and structure of the data, and does not provide support for predicates or projection. We perform a more comprehensive comparison against CRAM in §6.3.

3. PRINCIPLES FOR SCIENTIFIC ANALYSIS SYSTEMS

Most prior work on scientific computing has been focused on linear algebra and other problems that can be structured as a matrix or network. However, in several of the emerging data-driven scientific disciplines, data is less rigorously structured. As discussed in §2, scientists have been developing ad hoc solutions to process this data. In this section, we discuss the common characteristics of workloads in these emerging scientific areas. Given these characteristics, we describe a way to decompose data processing and storage systems so that they can efficiently implement important processing patterns *and* provide the necessary access methods for both computational researchers and field scientists.

3.1 Workloads

There are several common threads that unify the diverse set of applications that make up scientific computing. When looking at the data that is used in different fields, several trends pop out:

1. Scientific data tends to be rigorously associated with coordinates in some domain. These coordinate systems vary, but can include:

    - Time (e.g., fMRI data, particle simulations)
    - Chromosomal position (e.g., genomic read alignments and variants)
    - Position in space (imaging data, some sensor datasets)

2. For aggregated data, we may want to slice data into many different views. For example, for time domain data aggregated from many sensors, scientists may want to perform analyses by slicing across a single point in time, or by slicing across a single sensor. In genomics, we frequently aggregate data across many people from a given population. Once we have done this first aggregation, we may want to then slice the data by subsets of the population, or by regions of the genome (e.g., specific genes of interest).

There are two important consequences of the characteristics above. First, since data is attached to a coordinate system, the coordinate system itself may impose logical processing patterns. For example, for time domain data, we may frequently need to run functions that pass a sliding window, e.g., for convolution. Second, the slicing of aggregated data is frequently used to perform analyses across subsets of a larger dataset. This is common if we want to study a specific phenomenon, like the role of a gene in a disease (a common analysis in the TCGA [42]), or the measured activity in a single lobe of the brain while performing a task. Since the datasets we are processing are very large[1], it is often uneconomical to colocate data with processing nodes.

In this paper, we will tend to focus on compute tasks that are not simulation heavy; most simulation heavy tasks are dominated by communication between nodes due to the tight coupling of simulation elements. Because of this, they are not a good fit for the shared-nothing architectures common to "big data" systems. In non-simulation fields, the exact processing techniques and algorithms vary considerably, but common processing trends do exist:

1. There is increasing reliance on statistical methods. The `Thunder` pipeline makes heavy use of the MLI/MLLib statistical libraries [14, 38], and tools like the GATK perform multiple rounds of statistical refinement [11].

2. Data parallelism is very common. This varies across applications; in some applications (like genomics), we may leverage the independence of sites across a coordinate system and process individual coordinate regions in parallel. For other systems, we may have matrix

---

[1]For example, the Acute Myeloid Leukemia subset of the TCGA alone is over 4 TB in size.

calculations which can be parallelized [38], or we may be able to run processing in parallel across samples or traces.

Additionally, there are several different emerging use cases for scientific data processing and storage systems. These different use cases largely correspond to different points in the lifecycle of the data:

- **Batch processing:** After the initial acquisition of raw sensor data (e.g., raw DNA reads, brain electrode traces, telescope images), we use a batch processing pipeline (e.g., `Thunder` or the GATK) to perform some dimensionality reduction/statistical summarization of the data. This is generally used to extract notable features from the data, such as turning raw genomic reads into variant alleles, or identifying areas of activity in neuroscience traces. These tasks are unlikely to have any interactive component, and are likely to be long running compute jobs.

- **Ad hoc exploration:** Once the batch processing has completed, there is often a need for exploratory processing of the results. For example, when studying disease genetics, a geneticist may use the variant/genotype statistics to identify genomic sites with statistically significant links to the disease phenotype. Data exploration tasks have a significant user facing/interactive nature, and are generally performed by scientists who may be programming laypeople.

- **Data warehousing:** In large scientific projects, it is common to make data available to the members of the scientific community through some form of warehouse service (e.g., the Cancer Genomics Hub, CGHub, for the TCGA). As is the case for all data warehousing, this implies that queries must be made reasonably efficient, even though the data is expected to be cold. To reduce the cost of storing data, we may prioritize compression here; this has led to the creation of compressed storage formats like CRAM [15].

In this paper, we design a system which can achieve all of the above goals. The genomics and astronomy pipelines we demonstrate achieve improvements in batch processing performance, and allow for interactive/exploratory analysis through both Scala and Python. Through the layering principles we lay out in the next section and the performance optimizations we introduce in §5.2, we make our system useful for warehousing scientific data.

3.2 Layering

As discussed in the prior section, the processing patterns being applied to scientific data shift widely as the data itself ages. Because of this, we want to design a scientific data processing system that is flexible enough to accommodate our different use cases. At the same time, we want to ensure that the components in the system are well isolated so that we avoid bleeding functionality across the stack. If we bleed functionality across layers in the stack, we make it more difficult to adapt our stack to different applications.

Additionally, as we discuss in §4.1, improper separation of concerns can actually lead to errors in our application.

These concerns are very similar to the factors that led to the development of the Open Systems Interconnection (OSI) model and Internet Protocol (IP) stack for networking services [47]. The networking stack models were designed to allow the mixing and matching of different protocols, all of which existed at different functional levels. The success of the networking stack model can largely be attributed to the "narrow waist" of the stack, which simplified the integration of a new protocol or technology by ensuring that the protocol only needed to implement a single interface to be compatible with the rest of the stack.

Unlike conventional scientific systems which leverage custom data formats like BAM/SAM [24], or CRAM [15], we believe that the use of an explicit schema for data interchange is critical. In our stack model shown in Figure 1, the schema becomes the "narrow waist" of the stack. Most importantly, placing the schema as the narrow waist enforces a strict separation between data storage/access and data processing. Additionally, this enables literate programming techniques which can clarify the data model and access patterns.


Figure 1: A Stack Model for Scientific Computing

The seven layers of our stack model are decomposed as follows, traditionally numbered from bottom to top:

1. **Physical Storage:** This layer coordinates data writes to physical media.

2. **Data Distribution:** This layer manages access, replication, and distribution of the files that have been written to storage media.

3. **Materialized Data:** This layer encodes the patterns for how data is encoded and stored. This layer determines I/O bandwidth and compression.

4. **Data Schema:** This layer specifies the representation of data, and forms the narrow waist of the stack which separates access from execution.

5. **Evidence Access:** This layer provides us with primitives for processing data, and allows us to transform data into different views and traversals.

6. **Presentation:** This layer enhances the data schema with convenience methods for performing common tasks and accessing common derived fields from a single element.

7. **Application:** At this level, we can use our evidence access and presentation layers to compose the algorithms to perform our desired analysis.

A well defined software stack has several other significant advantages. By limiting application interactions with layers lower than the presentation layer, application developers are given a clear and consistent view of the data they are processing, and this view of the data is independent of whether the data is local or distributed across a cluster or cloud. By separating the API from the data access layer, we improve flexibility. With careful design in the data format and data access layers, we can seamlessly support conventional whole file access patterns, while also allowing easy access to small slices of files. By treating the compute substrate and storage as separate layers, we also drastically increase the portability of the APIs that we implement.

As we discuss in more detail in §4.1, current scientific systems bleed functionality between stack layers. An exemplar is the SAM/BAM and CRAM formats, which expect data to be sorted by genomic coordinate. This modifies the layout of data on disk (level 3, Materialized Data), which then constrains how applications traverse datasets (level 5, Evidence Access). Beyond constraining applications, this leads to bugs in applications which are difficult to detect.[2] To resolve this, we demonstrate several ways to efficiently implement conventional scientific traversals in §5. These traversals are implemented in the evidence access layer, and are independent of anything below the schema.

The idea of decomposing scientific applications into a stack model is not new; Bafna et al [5] made a similar suggestion in 2013. We borrow some vocabulary from Bafna et al, but our approach is differentiated in several critical ways:

- Bafna et al consider the stack model specifically in the context of data management systems for genomics; as a result, they bake current technologies and design patterns into the stack. In our opinion, a good stack design should serve to abstract layers from methodologies/implementations. If not, future technology trends

may obsolete a layer of the stack and render the stack irrelevant.

- Bafna et al define a binary data format as the narrow waist in their stack, instead of a schema. While these two seem interchangeable, they are not in practice. A schema is a higher level of abstraction and allows for data serialization techniques to be changed as long as the same schema is still provided.

- Notably, Bafna et al use this stack model to motivate GQL [19]. While a query system should provide a way to process and transform data, Bafna et al instead move this system down to the data materialization layer. We feel that this inverts the semantics that a user of the system would prefer and makes the system less general.

Our stack enables us to serve the use cases we have outlined in §3.1. By using Parquet as a storage format, we are able to process the data in many Hadoop-based systems. We implement high performance batch and interactive processing with Spark, and can delegate to systems like Impala and Spark-SQL for warehousing.

4. CASE STUDIES

To validate our architectural choices, we have implemented pipelines for processing short read genomic data and astronomy image processing. Both of these pipelines are implemented using Spark [46], Avro [2], and Parquet [3]. We have chosen these two applications as they fit in different areas in the design space. Specifically, the genomics pipeline makes heavy use of statistical processing techniques over semistructured data, while the astronomy application has a traditional matrix structure.

4.1 Genomics Pipeline

Contemporary genomics has been revolutionized by "next generation" sequencing technologies (NGS), which have driven a precipitous drop in the cost of running genomic assays [29]. Although there are a variety of sequencing technologies in use, the majority of sequence data comes from the Illumina sequencing platform, which uses a "sequencing-by-synthesis" technique to generate short read data [28]. Short read refers to sequencing run will generate many reads that are between 50 and 250 bases in length. In addition to adjusting the length of the reads, we can control the amount of the data that is generated by changing the amount of the genome that we sequence, or the amount of redundant sequencing that we perform (the average number of reads that covers each base, or *coverage*). A single human genome sequenced at $60\times$ coverage will produce approximately 1.4 billion reads, which is approximately 600 GB of raw data, or 225 GB of compressed data. For each read, we also are provided *quality scores*, which represent the likelihood that the base at a given position was observed.

One of the most common genomic analyses is *variant calling*, which is a statistical process to infer the sites at which a single individual varies from the reference genome.[3] To call

---

[2]The current best-practice implementations of the BQSR and Duplicate Marking algorithms both fail in certain corner-case alignments. These errors are caused because of the limitation on traversing reads in sorted order.

[3]The reference genome represents the "average" genome for a species. The Human Genome Project [21] assembled the first human reference genome.

variants, we perform the following steps:

1. **Alignment:** For each read, we find the position in the genome that the read is most likely to have come from. As an exact search is too expensive, there has been an extensive amount of research which has focused on indexing strategies for improving alignment performance [23, 25, 44]. This process is parallel per sequenced read.

2. **Pre-processing:** After reads have been aligned to the genome, we perform several preprocessing steps to eliminate systemic errors in the reads. This may involve recalibrating the observed quality scores for the bases, or locally optimizing the read alignments. We will present a description of several of these algorithms in §4.1; for a more detailed discussion, we refer readers to Appendix B, DePristo et al [11] and Massie et al [26].

3. **Variant calling:** Variant calling is a statistical process which uses the read alignments and the observed quality scores to compute whether a given sample matches or diverges from the reference genome. This process is typically parallel per position or region in the genome.

4. **Filtration:** After variants have been called, we want to filter out false positive variant calls. We may perform queries to look for variants with borderline likelihoods, or we may look for clusters of variants, which may indicate that a local error has occurred. This process may be parallel per position, may involve complex traversals of the genomic coordinate space, or may require us to fit a statistical model to all or part of the dataset.

This process is very expensive to run; the current best practice pipeline uses the BWA tool [23] for alignment and the GATK [27, 11] for pre-processing, variant calling, and filtration. Current benchmark suites have measured this pipeline as taking between 90 and 130 hours to run end-to-end [40]. Recent projects have achieved $5 - 10\times$ improvements in alignment and variant calling performance [44, 33], which makes the pre-processing stages the performance bottleneck. We have focused on implementing the four most time consuming pre-processing stages, as well as `flagstat`, a command which is used for quality control (QC). In the remainder of this section, we describe the stages that we have implemented, and the techniques we have used to improve performance and accuracy. We omit detailed pseudocode here, but make pseudocode available in Appendix B.

Sorting. This phase sorts all reads by the position of the start of their alignment. The implementation of this algorithm is trivial, as Spark provides a sort primitive [46]; we solely need to define an ordering for genomic coordinates, which is well defined. In practice, an explicit sort is unnecessary when using the rest of our MapReduce-based pipeline. We have included sort to enable the use of legacy tools which require sorted input.

Duplicate Removal. During the process of preparing DNA for sequencing, reads are duplicated by errors during the sample preparation and polymerase chain reaction (PCR) stages. Detection of duplicate reads requires matching all reads by their position and orientation after read alignment. Reads with identical position and orientation are assumed to be duplicates. When a group of duplicate reads is found, each read is scored, and all but the highest quality read are marked as duplicates.

We have validated our duplicate removal code against Picard [31], which is used by the GATK for Marking Duplicates. Our implementation is fully concordant with the Picard/GATK duplicate removal engine, with the exception of chimeric read pairs.[4] Specifically, because Picard's traversal engine is restricted to processing linearly sorted alignments, Picard mishandles these alignments. Since our engine is not constrained by the underlying layout of data on disk, we are able to properly handle chimeric read pairs.

Local Realignment. In local realignment, we correct areas where variant alleles cause reads to be locally misaligned from the reference genome.[5] In this algorithm, we first identify regions as targets for realignment. In the GATK, this is done by traversing sorted read alignments. In our implementation, we implement this as a fold over partitions where we generate targets, and then we merge the tree of targets. This allows us to eliminate the data shuffle needed to achieve the sorted ordering. As part of this fold, we must compute the convex hull of overlapping regions in parallel. We discuss this in more detail in §5.1

After we have generated the targets, we associate reads to an overlapping target (if one exists). After associating reads to realignment targets, we run the heuristic realignment algorithm which is described in Appendix **??**.

Base Quality Score Recalibration (BQSR). During the sequencing process, systemic errors occur that lead to the incorrect assignment of base quality scores. In this step, we label each base that we have sequenced with an *error covariate*. For each covariate, we count the total number of bases that we saw, as well as the total number of bases within the covariate that do not match the reference genome. From this data, we apply a Yates' correction to each covariate to recalculate the quality scores of all member bases:

$$P_{err}^{cov} = \frac{\text{\# errors} \in cov + 1}{\text{\# observations} \in cov + 2} \qquad (1)$$

We have validated the concordance of our BQSR implementation against the GATK. The overall differences are shown in Figure 2 (**FIXME: Frank + Jey to update this figure**). Across both tools, only 5000 of the $\sim$180B bases in the high-coverage NA12878 genome dataset differ. After in-

---

[4]In a chimeric read pair, the two reads in the read pairs align to different chromosomes; see Li et al [23].

[5]This is typically caused by the presence of insertion/deletion (INDEL) variants; see DePristo et al [11].

vestigating this discrepancy, we have determined that this is due to an error in the GATK, where paired-end reads are mishandled if the two reads in the pair overlap.


RMSE: 2.23
Exact Matches: 23.33%
Yates

Figure 2: BQSR Concordance vs. the GATK

For current implementations of these read processing steps, performance is limited by disk bandwidth [12]. This bottleneck exists because the operations read in a SAM/BAM file, perform a small amount of processing, and write the data to disk as a new SAM/BAM file. We achieve a performance bump by performing our processing iteratively in memory. The four read processing stages can then be chained together, eliminating three long writes to disk and an additional three long reads from disk. Additionally, by rethinking the design of our algorithms, we are able to reduce overhead in several other ways:

1. Current algorithms require the reference genome to be present on all nodes. This is then used to look up the reference sequence that overlaps all reads. The reference genome is several gigabytes in size, and performing a lookup in the reference genome can be costly due to its size. Instead, we leverage a field in our schema to embed information about the reference in each read. This allows us to avoid broadcasting the reference, and provides O(1) lookup.

2. Shared-memory genomics applications tend to be impacted significantly by false sharing of data structures [44]. Instead of having data structures that are modified in parallel, we restructure our algorithms so that we only touch data structures from a single thread, and then merge structures in a reduce phase. This improves the performance of covariate calculation during BQSR and the target generation phase of local realignment.

3. In a naïve implementation, the local realignment and duplicate marking tasks can suffer from bad stragglers. This occurs due to a large amount of reads that either do not associate to a realignment target, or that are unaligned. We pay special attention to these cases by manually randomizing the partitioning for these reads. This resolves load imbalance and mitigates stragglers.

4. For the Flagstat command, we are able to project a limited subset of fields. Flagstat touches fewer than 10 fields, which account for less than 3% (**FIXME: Frank! check this number!**) of space on disk. We discuss the performance implications of this further in §6.3.

These techniques allow us to achieve a $> 50\times$ performance improvement over current tools, and scalability beyond 128 machines. We perform a detailed performance review in §6.1

4.2 Astronomy Image Processing

The Montage [18] application is a typical astronomy image processing pipeline that builds mosaics from small image tiles obtained from telescopes with the requirement of preserving the energy quantity and position for each pixel between the input and output images. The pipeline has multiple stages that can be grouped into the following phases.

Tile Reprojection. The raw input images are reprojected with the defined scale as required in the final mosaic.

Background Modeling. This phase smoothes out the background levels between each pair of overlapped images and fits a plane to each of them. The phase can be further divided into steps of overlap calculation, difference image creation, and plane-fitting coefficient calculation.

Background Matching. The background matching phase applies background removal to the reprojected images with the best solution derived from previous phase to smooth out the overlap regions.

Tile Mosaicing. The tile masoning phase coadds all corrected images (after applying background matching to the reprojected images) into a aggregated mosaic file. This phase also involves a metadata processing stage before the mosaicing.

The tile reproduction, background modeling, and background matching phases are embarrassingly parallel, with each task (both computation and I/O) running independent from other tasks in the same phase. The tile mosaicing phase aligns the tiles of the matrix and applies a function (the function can be average, median, or count) to the overlapped pixels. All functions can be applied to the overlapped pixels in an embarrassingly parallel manner, while the I/O are done through a single process.

We care specifically about the tile mosaicing phase as the

current implementation requires a preceding stage to summarize the metadata of all corrected images to produce a metadata table containing the tile positioning information. This particular stage has to read all input files into memory, but only accesses a small portion of the file. Also the current implementation parallelize the computation with each matrix row as the element, which results in inefficient input images replication when executed in a distributed environment. We address the first by explicitly declaring the image data schema and storing the images in a columnar store, so that all input images can be loaded to memory just once with all subsequent computation in memory. The metadata processing can access data that is continuous on a disk, while the rest of the I/O can be done in parallel.

5. DATA ACCESS OPTIMIZATIONS FOR SCIENTIFIC PROCESSING

In §3.1, we discussed several processing patterns which were important for scientific data processing. In this section, we introduce optimizations for two important use cases. First, we present a join pattern which enables processing that traverses a coordinate system. This enables distributed implementations of important genomics algorithms. Second, we implement an efficient method for applying predicates and projections into data which is stored in a remote block store. This allows us to improve the efficiency of accessing remotely staged data.

5.1 Coordinate System Joins

Genome informatics encompasses a wide array of experimental techniques and platforms, but almost all of these methods share one characteristic in common: they produce datapoints which are tied to locations in the genome through the use of genomic coordinates. The actual genome inside each cell is a mammoth molecule, a collection of DNA polymers coated with (and wrapped around) proteins and packed into the nucleus in a complex 3-dimensional shape. Bioinformaticians avoid this complexity by representing all 3.2 billion bases of the human genome as a single long string (of A's, T', G's, and C's); the "tying" of a datapoint or observation to the genome is represented by associating the data with a 1 dimensional point or interval. The reference genome defines a 1-dimensional coordinate space against which all genomic phenomena are marked and measured.

A platform for scientific data processing in genomics needs to understand these 1-dimensional coordinate systems because these become the basis on which data processing is parallelized. For example, when calling variants from sequencing data, the sequence data which is localized to a single genomic region (or "locus") can be processed independently from the data localized to a different region, as long as the regions are far enough apart.

Not only parallelization, but many of the core algorithms and methods for data aggregation in genomics are phrased in terms of geometric primitives on 1-D intervals and points: distance, overlap, and containment. An algorithm for calculating quality control metrics may try to calculate "coverage," a count of how many reads overlap each base in the genome. A method for filtering and annotating potential variants might assess the validity of a variant by the quality and mapping characteristics of all reads which overlap the putative variant.

In order to support these algorithms, we have implemented a "region" or "spatial" join method, whose outline in code is given below. The region join operation takes as input two sets (RDDs) of ReferenceRegions, a data structure that represents intervals along the 1-D genomics coordinate space. It produces, as output, the set of all ReferenceRegion pairs (one from each of the two input sets) such that the two regions overlap each other in terms of genomics spatial coordinates.

```
def partitionAndJoin[T, U](sc: SparkContext,
  baseRDD: RDD[T],
  joinedRDD: RDD[U],
  tMapping: ReferenceMapping[T],
  uMapping: ReferenceMapping[U]): RDD[(T, U)] = {

  val collectedLeft = baseRDD
      .map(t => (tMapping.getReferenceName(t),
                   tMapping.getReferenceRegion(t)))
      .groupBy(_._1)
      .map(t => (t._1, t._2.map(_._2)))
      .collect()
      .toSeq

  val multiNonOverlapping =
      new MultiContigNonoverlappingRegions(collectedLeft)

  val regions = sc.broadcast(multiNonOverlapping)

  val smallerKeyed: RDD[(ReferenceRegion, T)] =
    baseRDD.keyBy(t => regions.value.regionsFor(t).head)

  val largerKeyed: RDD[(ReferenceRegion, U)] =
    joinedRDD.filter(regions.value.filter(_))
      .flatMap(t => regions.value.regionsFor(t)
                   .map((r: ReferenceRegion) => (r, t)))

  val joined: RDD[(ReferenceRegion, (T, U))] =
    smallerKeyed.join(largerKeyed)

  val filtered: RDD[(ReferenceRegion, (T, U))] =
    joined.filter({
      case (rr: ReferenceRegion, (t: T, u: U)) =>
  tMapping.getReferenceRegion(t)
    .overlaps(uMapping.getReferenceRegion(u))
})

  filtered.map(rrtu => rrtu._2)
}
```

FIXME: To find the non-overlapping regions...

5.2 Loading Remote Data

Another challenge faced by scientific systems is where to store the initial data files and how to load them efficiently. Today, Spark is usually run in conjunction with the HDFS portion of the Hadoop stack—HDFS provides data locality, access to local disk on each node of the Spark cluster, and robustness to node failure. However, HDFS imposes significant constraints on running a Spark system in virtualized or commodity computing (e.g. "cloud") environments. It is easy to scale an HDFS-based system up to larger numbers of nodes, but harder to remove nodes when the capacity is no longer needed.

If we are willing to forgo the advantages of local disk and data locality provided by HDFS, however, we may be able to relax some of these other restrictions and build a Spark-based cluster whose size is more easily adjusted to the changing demands of the computation. By storing our data in higher-latency, durable, cheaper block storage (e.g. S3) we can also exploit the varying requirements of data availability—not all datasets need to be kept "hot" in HDFS at all times, but can be accessed in a piecemeal or parallelized manner through S3 interfaces.

Spark provides a particularly convenient abstraction for writing these new data access methods. By implementing our own data-loading RDD, we are able to allow a Spark cluster to access Parquet files stored in S3 in parallel (each partition in the RDD reflects a row group in the corresponding Parquet file). For Parquet files containing records that reflect known genomics datatypes (that are mapped to genomic locations, for example) we generate simple index files for each Parquet file. Each index file lists the complete set of row groups for the Parquet file, as well which genomic regions contain data points within each row group. Our parallelized data loader reads this index file and restricts the partitions in the data loading RDD it creates to only those Parquet row groups which possibly contain data relevant to the user's query.

6. PERFORMANCE

In this paper, we have discussed ways to improve the performance of scientific workloads which are being run on commodity MapReduce systems by rethinking how we decompose and build algorithms. In this section, we review the improvements in performance that we are able to unlock. We achieve near-linear speedup across 128 nodes for a genomics workload, and achieve a 3× performance improvement over the current best MPI-based system for the Montage astronomy application. Additionally, we achieve 30-50% compression over current file formats when storing to disk.

6.1 Genomics Workloads

FIXME: These are old/outdated/need to be anonymized! Frank to update. Also, comment about data conversion cost.

Table 1 previews the performance of ADAM for *Sort*, *Mark Duplicates*, and *Flagstat*. The tests in this table are run on the high coverage *NA12878* full genome BAM file that is available from the 1000 Genomes project; the HG00096 low coverage BAM from 1000 Genomes is used later in this section[6]. These tests have been run on the EC2 cloud, using the instance types listed. We compute the cost of running each experiment by multiplying the number of instances used by the total wall time for the run and by the cost of running a single instance of that type for an hour, which is the process Amazon uses to charge customers.

Table 2 describes the instance types. Memory capacity is

---

[6]The files used for these experiments can be found on the 1000 Genomes ftp site, `ftp.1000genomes.ebi.ac.uk` in directory `/vol1/ftp/data/NA12878/high_coverage_alignment/` for NA12878, and in directory `/vol1/ftp/data/HG00096/alignment/` for HG00096.

Table 1: *Sort*, *Mark Duplicates*, **and** *Flagstat* **Performance on NA12878**

Sort				
Software	**EC2 profile**	**Wall Time**	**Speedup**	**Cost**
Picard 1.103	1 `hs1.8xlarge`	17h 44m	1×	$81.57
ADAM 0.5.0	1 `hs1.8xlarge`	8h 56m	2×	$41.09
ADAM 0.5.0	32 `cr1.8xlarge`	33m	32×	$61.60
ADAM 0.5.0	100 `m2.4xlarge`	21m	52×	$56.00
Mark Duplicates				
Software	**EC2 profile**	**Wall Time**	**Speedup**	**Cost**
Picard 1.103	1 `hs1.8xlarge`	20h 22m	1×	$93.68
ADAM 0.5.0	100 `m2.4xlarge`	29m	42×	$79.26
Flagstat				
Software	**EC2 profile**	**Wall Time**	**Speedup**	**Cost**
SAMtools 0.1.19	1 `hs1.8xlarge`	25m 24s	1×	$1.95
ADAM 0.5.0	32 `cr1.8xlarge`	0m 46s	33×	$1.43

reported in Gibibytes (GiB), where 1 GiB is equal to 2^{30} bytes. Storage capacities are not reported in this table because disk capacity does not impact performance, but the number and type of storage drives is reported because aggregate disk bandwidth does impact performance. In our tests, the `hs1.8xlarge` instance is chosen to represent a workstation. Network bandwidth is constant across all instances.

Table 2: AWS Machine Types

Machine	Cost	Description
`hs1.8xlarge`	$4.60/hr/machine	16 cores, 117GiB RAM, 24× HDD
`cr1.8xlarge`	$3.50/hr/machine	32 cores, 244GiB RAM, 2× SDD
`m2.4xlarge`	$1.64/hr/machine	8 cores, 68.4GiB RAM, 2× HDD

As can be seen from these results, the ADAM pipeline is approximately twice as fast as current pipelines when running on a single node. Additionally, ADAM achieves speedup that is close to linear. This point is not clear from Table 1, as we change instance types when also changing the number of instances used. To clarify, Figure 3 presents speedup plots for the NA12878 high coverage genome and the HG00096 low coverage genome (16 GB BAM). These speedup measurements are taken on a dedicated cluster of 82 machines where each machine has 24 Xeon cores, 128GB of RAM, and 12 disks.

NA12878 sees linear speedup for both *Sort* and *Mark Duplicates* through 82 nodes. With 82 nodes, it takes 8.8 minutes to sort reads, and 14 minutes to mark duplicate reads across the 250GB file. HG00096 is a significantly smaller genome at 16 GB. Although the speedup from sorting diminishes after 16 nodes, duplicate marking sees speedup through 64 nodes. With HG00096 on 64 nodes, each machine in the cluster is responsible for processing only 250MB of data. Sorting completes in 3.3 minutes, and duplicate marking completes in 2.3 minutes. The speedup is limited by several factors:

- Although columnar stores have very high read performance, they are limited by write performance. Our tests exaggerate this penalty—as a variant calling pipeline will consume a large read file, but then output a variant call file that is approximately two orders of magnitude smaller, the write penalty will be reduced.

Figure 3: Speedup when running *Sort* and *Mark Duplicates* on NA12878 and HG00096

- Additionally, for large clusters, straggler elimination is an issue. Phases of both *Sort* and *MarkDuplicates* currently suffer from stragglers—we are in the process of addressing these issues.

However, as noted above, speedup continues until we reach approximately 1GB of data per node for sorting, or 250MB of data per node for duplicate marking. For a high coverage whole genome like NA12878, this should theoretically allow speedup through 250 nodes for sorting and 1,000 nodes for duplicate marking.

6.2 Astronomy Workloads

We use the 2MASS data and the Montage test case of 3x3 degree mosaicing with Galaxy m101 as the center. The tile mosaicing phase has 1.5 GB input data and produces a 1.2 GB aggregated output file. We compare the SparkmAdd performance against the MPI based parallel implementation in Montage v3.3 (referred as MPImAdd in the following text). The performance is measured on 1, 4, and 16 Amazon `m2.4xlarge` instances respectively. We use OrangeFS v2.8.8, a successor of PVFS [7], as the shared file system across the instances to support MPImAdd. All eight cores on each instance are used for both SparkmAdd and MPImAdd.


Figure 4: Speedup when running *mAdd* using MPI and Spark

As shown in Figure 4, the SparkmAdd runs 1.3x, 2.8x, 3.3x faster than MPImAdd on 1, 4, and 16 instances with a data compression rate of 2.8x for input and 1.4x for output. The performance improvement is contributed by multiple factors: less I/O amount, less I/O contention, and better data locality. We manage to combine the metadata processing stage with mAdd together, since we explicitly define the data schema in SparkmAdd, so that the same set of inputs can just be read once. We also configure Parquet to compress the input and output data, so the actual I/O amount is reduced. Parquet writes output files into HDFS in a contention free manner whereas MPI coordinates all processes to write to a single file in the shared file system. The Spark framework lets the computation benefit from data locality, while MPI distributes the computation across the available resources without concerning data placement.

6.3 Column Store Performance

FIXME: This uses old data and also needs to be anonymized. Frank to fix. Also, discuss CRAM, and columnar representations for in-memory data.

The decision to use a columnar store was based on the observation that most genomics applications are read-heavy. In this section, we aim to quantify the improvements in read performance that we achieve by using a column store, and how far these gains can be maximized through the use of projections.

Read Length and Quality:. This microbenchmark scans all the reads in the file, and collects statistics about the length of the read, and the mapping quality score. Figure 5 shows the results of this benchmark. Ideal speedup curves are plotted along with the speedup measurements. The speedups in this graph are plotted relative to the single machine Hadoop-BAM run.

Figure 5: Read Length and Quality Speedup for HG00096

This benchmark demonstrates several things:

- ADAM demonstrates superior scalability over Hadoop-BAM. This performance is likely because ADAM files can be more evenly distributed across machines. Hadoop-BAM must guess the proper partitioning when splitting files [30]. The Parquet file format that ADAM is based on is partitioned evenly when it is written[3]. This paritioning is critical for small files such as the one used in this test — in a 16 machine cluster, there is only 1 GB of data per node, so imbalances on the order of the Hadoop file split size (128 MB) can lead to 40% differences in loading between nodes.

- Projections can significantly accelerate computation, if the calculation does not need all the data. By projecting the minimum dataset necessary to calculate the target of this benchmark, we can accelerate the benchmark by 8×. Even a more modest projection leads to a 2× improvement in performance.

6.3.1 Predicate Pushdown

Predicate pushdown is one of the significant advantages afforded to us through the use of Parquet [3]. In traditional pipelines, the full record is deserialized from disk, and the filter is implemented as a Map/Reduce processing stage. Parquet provides us with predicate pushdown: we deserialize only the columns needed to identify whether the record should be fully deserialized. We then process those fields; if they pass the predicate, we then deserialize the full record.

We can prove that this process requires no additional reads. Intuitively, this is true as the initial columns read would need to be read if they were to be used in the filter later. However, this can be formalized. We provide a proof for this in Section C of the Appendix. In this section, we seek to provide an intuition for how predicate pushdown can improve the performance of actual genomics workloads. We apply predicate pushdown to implement read quality predication, mapping score predication, and predication by chromosome.

Quality Flags Predicate:. This microbenchmark scans the reads in a file, and filters out reads that do not meet a specified quality threshold. We use the following quality threshold:

- The read is mapped.

- The read is at its primary alignment position.

- The read did not fail vendor quality checks.

- The read has not been marked as a duplicate.

This threshold is typically applied as one of the first operations in any read processing pipeline.

Gapped Predicate:. In many applications, the application seeks to look at a region of the genome (e.g. a single chromosome or gene). We include this predicate as a proxy: the gapped predicate looks for 1 out of every n reads. This selectivity achieves performance analogous to filtering on a single gene, but provides an upper bound as we pay a penalty due to our predicate hits being randomly distributed.

To validate the performance of predicate pushdown, we ran the predicates described above on the HG00096 genome on a single AWS $m2.4xlarge$ instance. The goal of this was to determine the rough overhead of predicate projection.

Figure 6 shows performance for the gapped predicate sweeping n.

Figure 6: Gapped Predicate on HG00096

There are several conclusions to be drawn from this experiment:

- As is demonstrated by equation (**??**) in Section C of the Appendix, we see the largest speedup when our predicate read set is significantly smaller than our projection read set. For smaller projections, we see a fairly substantial reduction in speedup.

- As the number of records that passes the predicate drops, the number of reads done to evaluate the predicate begins to become a dominant term. This insight has one important implication: for applications that plan to access a very small segment of the data, accessing the data as a flat file with a predicate is likely not the most efficient access pattern. Rather, if the predicate is regular, it is likely better to access the data through a indexed database.

We additionally calculated the performance improvement that was achieved by using predicate pushdown instead of a Spark filter. Table 3 lists these results.

Table 3: Predicate Pushdown Speedup vs. Filtering

Predicate	Min	Avg	Max	Filtered
Locus	1.19	1.17	0.94	3%
Gapped, $n = 1$	1.0	1.0	1.01	0%
Gapped, $n = 2$	1.22	1.27	1.29	50%
Gapped, $n = 4$	1.31	1.42	1.49	75%
Gapped, $n = 8$	1.37	1.58	1.67	87.5%

These results are promising, as they indicate that predicate pushdown is faster than a Spark filter in most cases. We believe that the only case that showed a slowdown (max projection on Locus predicate) is due to an issue in our test setup. We therefore can conclude that filtering operations that can be performed statically at file load-in should be performed using predicate pushdown.

7. FUTURE WORK

This work leverages columnar storage to improve performance and compression of data on disk, with special emphasis on repetitive fields which can be run length encoded (RLE). While this improves disk performance, it has the side effect of making data consume significantly more space in memory than on disk. We are currently investigating techniques that leverage the immutability of data in our applications to reduce memory consumption. This may involve a pseudocolumnar data representation in memory, or changes to Parquet and Avro's deserialization codec which would reuse allocated objects.

It is worth noting that there are many significant scientific applications (such as genome assembly) that are expressed as traversal over graphs. Recent work by Simpson et al (ABySS, [36]) and Georganas et al [16] has focused on using MPI or Unified Parallel C (UPC) to roll their own distributed graph traversal. Both systems find that synchronization via message passing is a significant cost; specifically, the ABySS assembler experiences scaling problems because it thrashes portions of the graph across nodes during traversal. By building our system using Spark, we are able to leverage the GraphX processing library [43, 17]. We are in the process of developing a genome assembler using this library system, and believe that we can achieve improved performance through careful graph partitioning. This involves algorithmic changes to the graph creation and traversal phases to bypass "knotted" sections of the graph which

correspond to highly repetitive areas of the genome, which cause the major performance issues in MPI based assemblers.

8. CONCLUSION

In this paper, we have advocated for a clean architecture for decomposing components in a scientific system, and then demonstrated how to efficiently implement genomic and astronomy processing pipelines using the open source Avro, Parquet, and Spark systems [2, 3, 46]. We demonstrate > $50\times$ performance improvements over conventional genomics processing systems, along with linear strong scaling and a $2\times$ cost improvement. On the astronomy workload, we achieve a $3\times$ speedup over the current best MPI-based solution. Additionally, while an obvious advantage of using commodity systems is code reuse, we have gained several valuable features for free:

1. Beyond efficient parallel performance, Parquet is also supported as a data format by several database-like systems, like Spark-SQL and Impala. This allows us to support efficient database style processing, without needing to manually retrofit a tool like GQL to the data [19].

2. While the native Spark Scala API provides rich programming abstractions, scientists may prefer other language environments, like Python or R. Other scientific systems like SciDB [6] support native Python and R bindings. Likewise, we provide distributed processing using Python through PySpark.

By rethinking the architecture of scientific data management systems, we have been able to achieve significant scalability improvements while also expanding the methods which can be used to process the data. The architecture we have developed is clean and principled and minimizes the mixture-of-concerns which occurs in current scientific systems. Additionally, by applying our techniques to both astronomy and genomics, we have demonstrated that the techniques are applicable to both traditional matrix-based scientific computing, as well as novel scientific areas which have less structured data.

APPENDIX

A. REFERENCES

[1] D. Abadi, S. Madden, and M. Ferreira. Integrating compression and execution in column-oriented database systems. In *Proceedings of the 2006 ACM SIGMOD international conference on Management of Data (SIGMOD '06)*, pages 671–682. ACM, 2006.

[2] Apache. Avro. `http://avro.apache.org`.

[3] Apache. Parquet. `http://parquet.incubator.apache.org`.

[4] G. A. Auwera, M. O. Carneiro, C. Hartl, R. Poplin, G. del Angel, A. Levy-Moonshine, T. Jordan, K. Shakir, D. Roazen, J. Thibault, et al. From FastQ data to high-confidence variant calls: The Genome Analysis Toolkit best practices pipeline. *Current Protocols in Bioinformatics*, pages 11–10, 2013.

[5] V. Bafna, A. Deutsch, A. Heiberg, C. Kozanitis, L. Ohno-Machado, and G. Varghese. Abstractions for genomics. *Communications of the ACM*, 56(1):83–93, 2013.

[6] P. G. Brown. Overview of SciDB: large scale array storage, processing and analysis. In *Proceedings of the 2010 ACM SIGMOD International Conference on Management of Data (SIGMOD '10)*, pages 963–968. ACM, 2010.

[7] P. H. Carns, W. B. Ligon, III, R. B. Ross, and R. Thakur. PVFS: a parallel file system for linux clusters. In *Proceedings of the 4th annual Linux Showcase & Conference*, pages 28–28. USENIX Association, 2000.

[8] J. P. Cunningham. Analyzing neural data at huge scale. *Nature Methods*, 11(9):911–912, 2014.

[9] J. Dean and S. Ghemawat. MapReduce: simplified data processing on large clusters. In *Proceedings of the 6th Symposium on Operating System Design and Implementation (OSDI '04)*. ACM, 2004.

[10] J. Dean and S. Ghemawat. MapReduce: simplified data processing on large clusters. *Communications of the ACM*, 51(1):107–113, 2008.

[11] M. A. DePristo, E. Banks, R. Poplin, K. V. Garimella, J. R. Maguire, C. Hartl, A. A. Philippakis, G. del Angel, M. A. Rivas, M. Hanna, et al. A framework for variation discovery and genotyping using next-generation DNA sequencing data. *Nature Genetics*, 43(5):491–498, 2011.

[12] Y. Diao, A. Roy, and T. Bloom. Building highly-optimized, low-latency pipelines for genomic data analysis. In *Proceedings of the 7th Conference on Innovative Data Systems Research (CIDR '15)*, 2015.

[13] G. England. 100,000 genomes project. `https://www.genomicsengland.co.uk/`.

[14] J. Freeman, N. Vladimirov, T. Kawashima, Y. Mu, N. J. Sofroniew, D. V. Bennett, J. Rosen, C.-T. Yang, L. L. Looger, and M. B. Ahrens. Mapping brain activity at scale with cluster computing. *Nature Methods*, 11(9):941–950, 2014.

[15] M. H.-Y. Fritz, R. Leinonen, G. Cochrane, and E. Birney. Efficient storage of high throughput DNA sequencing data using reference-based compression. *Genome Research*, 21(5):734–740, 2011.

[16] E. Georganas, A. Buluç, J. Chapman, L. Oliker, D. Rokhsar, and K. Yelick. Parallel de bruijn graph construction and traversal for de novo genome assembly. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC Õ14)*, 2014.

[17] J. E. Gonzalez, R. S. Xin, A. Dave, D. Crankshaw, M. J. Franklin, and I. Stoica. GraphX: Graph processing in a distributed dataflow framework. In *Proceedings of the 16th Symposium on Operating System Design and Implementation (OSDI '14)*. ACM, 2014.

[18] J. C. Jacob, D. S. Katz, G. B. Berriman, J. C. Good, A. C. Laity, E. Deelman, C. Kesselman, G. Singh, M.-H. Su, T. A. Prince, and R. Williams. Montage: a grid portal and software toolkit for science-grade astronomical image mosaicking. *International Journal of Computational Science and Engineering*, 4(2):73–87, 2009.

[19] C. Kozanitis, A. Heiberg, G. Varghese, and V. Bafna. Using Genome Query Language to uncover genetic variation. *Bioinformatics*, 30(1):1–8, 2014.

[20] A. Lamb, M. Fuller, R. Varadarajan, N. Tran, B. Vandiver, L. Doshi, and C. Bear. The Vertica analytic database: C-store 7 years later. *Proceedings of the VLDB Endowment*, 5(12):1790–1801, 2012.

[21] E. S. Lander, L. M. Linton, B. Birren, C. Nusbaum, M. C. Zody, J. Baldwin, K. Devon, K. Dewar, M. Doyle, W. FitzHugh, et al. Initial sequencing and analysis of the human genome. *Nature*, 409(6822):860–921, 2001.

[22] B. Langmead, M. C. Schatz, J. Lin, M. Pop, and S. L. Salzberg. Searching for SNPs with cloud computing. *Genome Biology*, 10(11):R134, 2009.

[23] H. Li and R. Durbin. Fast and accurate long-read alignment with Burrows–Wheeler transform. *Bioinformatics*, 26(5):589–595, 2010.

[24] H. Li, B. Handsaker, A. Wysoker, T. Fennell, J. Ruan, N. Homer, G. Marth, G. Abecasis, R. Durbin, et al. The sequence alignment/map format and SAMtools. *Bioinformatics*, 25(16):2078–2079, 2009.

[25] Y. Li, A. Terrell, and J. M. Patel. WHAM: A high-throughput sequence alignment method. In *Proceedings of the 2011 ACM SIGMOD International Conference on Management of Data (SIGMOD '11)*, SIGMOD '11, pages 445–456, New York, NY, USA, 2011. ACM.

[26] M. Massie, F. Nothaft, C. Hartl, C. Kozanitis, A. Schumacher, A. D. Joseph, and D. A. Patterson. ADAM: Genomics formats and processing patterns for cloud scale computing. Technical report, UCB/EECS-2013-207, EECS Department, University of California, Berkeley, 2013.

[27] A. McKenna, M. Hanna, E. Banks, A. Sivachenko, K. Cibulskis, A. Kernytsky, K. Garimella, D. Altshuler, S. Gabriel, M. Daly, et al. The Genome Analysis Toolkit: a MapReduce framework for analyzing next-generation DNA sequencing data. *Genome Research*, 20(9):1297–1303, 2010.

[28] M. L. Metzker. Sequencing technologies—the next generation. *Nature Reviews Genetics*, 11(1):31–46, 2009.

[29] NHGRI. DNA sequencing costs.

http://www.genome.gov/sequencingcosts/.

[30] M. Niemenmaa, A. Kallio, A. Schumacher, P. Klemelä, E. Korpelainen, and K. Heljanko. Hadoop-BAM: directly manipulating next generation sequencing data in the cloud. *Bioinformatics*, 28(6):876–877, 2012.

[31] T. B. I. of Harvard and MIT. Picard. http://broadinstitute.github.io/picard/, 2014.

[32] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, et al. Scikit-learn: Machine learning in python. *The Journal of Machine Learning Research*, 12:2825–2830, 2011.

[33] A. Rimmer, H. Phan, I. Mathieson, Z. Iqbal, S. R. Twigg, A. O. Wilkie, G. McVean, G. Lunter, W. Consortium, et al. Integrating mapping-, assembly-and haplotype-based approaches for calling variants in clinical sequencing applications. *Nature Genetics*, 46(8):912–918, 2014.

[34] E. E. Schadt, M. D. Linderman, J. Sorenson, L. Lee, and G. P. Nolan. Computational solutions to large-scale data management and analysis. *Nature Reviews Genetics*, 11(9):647–657, 2010.

[35] M. C. Schatz. CloudBurst: highly sensitive read mapping with MapReduce. *Bioinformatics*, 25(11):1363–1369, 2009.

[36] J. T. Simpson, K. Wong, S. D. Jackman, J. E. Schein, S. J. Jones, and I. Birol. ABySS: a parallel assembler for short read sequence data. *Genome Research*, 19(6):1117–1123, 2009.

[37] N. Siva. 1000 genomes project. *Nature Biotechnology*, 26(3):256–256, 2008.

[38] E. R. Sparks, A. Talwalkar, V. Smith, J. Kottalam, X. Pan, J. Gonzalez, M. J. Franklin, M. I. Jordan, and T. Kraska. MLI: An API for distributed machine learning. In *13th IEEE International Conference on Data Mining (ICDM '13)*, pages 1187–1192. IEEE, 2013.

[39] L. D. Stein et al. The case for cloud computing in genome informatics. *Genome Biology*, 11(5):207, 2010.

[40] A. Talwalkar, J. Liptrap, J. Newcomb, C. Hartl, J. Terhorst, K. Curtis, M. Bresler, Y. S. Song, M. I. Jordan, and D. Patterson. SMASH: A benchmarking toolkit for human genome variant calling. *Bioinformatics*, page btu345, 2014.

[41] M. J. Turk, B. D. Smith, J. S. Oishi, S. Skory, S. W. Skillman, T. Abel, and M. L. Norman. yt: A multi-code analysis toolkit for astrophysical simulation data. *The Astrophysical Journal Supplement Series*, 192(1):9, 2011.

[42] J. N. Weinstein, E. A. Collisson, G. B. Mills, K. R. M. Shaw, B. A. Ozenberger, K. Ellrott, I. Shmulevich, C. Sander, J. M. Stuart, C. G. A. R. Network, et al. The cancer genome atlas pan-cancer analysis project. *Nature Genetics*, 45(10):1113–1120, 2013.

[43] R. S. Xin, J. E. Gonzalez, M. J. Franklin, and I. Stoica. GraphX: A resilient distributed graph system on Spark. In *First International Workshop on Graph Data Management Experiences and Systems*, page 2. ACM, 2013.

[44] M. Zaharia, W. J. Bolosky, K. Curtis, A. Fox, D. Patterson, S. Shenker, I. Stoica, R. M. Karp, and T. Sittler. Faster and more accurate sequence

alignment with SNAP. *arXiv preprint arXiv:1111.5572*, 2011.

[45] M. Zaharia, M. Chowdhury, T. Das, A. Dave, J. Ma, M. McCauley, M. Franklin, S. Shenker, and I. Stoica. Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing. In *Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation (NSDI '12)*, page 2. USENIX Association, 2012.

[46] M. Zaharia, M. Chowdhury, M. J. Franklin, S. Shenker, and I. Stoica. Spark: cluster computing with working sets. In *Proceedings of the 2nd USENIX conference on Hot topics in Cloud Computing (HotCloud '10)*, page 10, 2010.

[47] H. Zimmermann. OSI reference model–the ISO model of architecture for open systems interconnection. *IEEE Transactions on Communications*, 28(4):425–432, 1980.

B. GENOMICS PSEUDOCODE

With the exception of the `Sort` stage, the genomics preprocessing stages seek to collect global information that is not available to the sequencer or aligner to correct sequenced or aligned data. We include pseudocode for the `Mark Duplicates`, `BQSR`, and `Indel Realignment` stages below.

B.1 BQSR

Base Quality Score Recalibration is an important data normalization step in the variant calling pipeline, and is the second costliest step after realignment. BQSR is composed of two stages:

1. We perform a parallel aggregate to compute the statistics for all covariate buckets. We put all bases into their appropriate covariate and then identify how many bases in each covariate mismatch against the reference at their aligned position.

2. Once the covariate statistics have been computed, we rewrite the base quality scores per base. This can be done through a map stage.

The covariates capture characteristics of the sequenced base that are expected to be correlated with error. We support the following covariates:

- `Measured Score`: This is the quality score originally measured by the sequencer.

- `Cycle`: This is the cycle the base was sequenced in. Errors are known to be most common at the start and end of reads [11].

- `Dinucleotide`: This covariate tracks the value of this base and the nucleotide that immediately preceded it. Sequencers struggle with reading certain patterns of bases.

We compute the covariate table with the following code:

```
val dataset: RDD[(CovariateKey, Residue)] = {
  def shouldIncludeRead(read: DecadentRead) =
    read.isCanonicalRecord &&
```

```
    read.alignmentQuality.exists(_ > QualityScore.zero)
    && read.passedQualityChecks

  def shouldIncludeResidue(residue: Residue) =
    residue.quality > QualityScore.zero &&
      residue.isRegularBase &&
      !residue.isInsertion &&
      !knownSnps.value.isMasked(residue)

  def observe(read: DecadentRead): Seq[(CovariateKey,
    Residue)] =
    covariates(read).zip(read.residues).
      filter { case (key, residue) =>
        shouldIncludeResidue(residue) }

  input.filter(shouldIncludeRead).flatMap(observe)
}

val observed: ObservationTable = {
  dataset.
    map { case (key, residue) => (key,
      Observation(residue.isSNP)) }.
    aggregate(ObservationAccumulator(covariates))(_ += _,
      _ ++= _).result
}
```

Once the table is recalibrated, we map the following function over the reads to write the recalibrated scores:

```
def recalibrate(read: DecadentRead): AlignmentRecord = {
  val record: AlignmentRecord = read.record
  AlignmentRecord.newBuilder(record).
    setQual(QualityScore.toString(computeQual(read))).
    setOrigQual(record.getQual).
    build()
}

def computeQual(read: DecadentRead): Seq[QualityScore] = {
  val origQuals = read.residues.map(_.quality)
  val newQuals = table(read)
  origQuals.zip(newQuals).map {
    case (origQ, newQ) =>
      // Keep original quality score if below
      // recalibration threshold
      if (origQ >= minAcceptableQuality) newQ else origQ
  }
}
```

B.2 Indel Realignment

Indel realignment is a two step process. First we identify regions that have evidence of an insertion or deletion. Then we realign reads to candidate haplotypes to minimize the overall weight of mismatches in the region. The quality of mismatches near an INDELs serves as a good proxy for the local quality of an alignment. This is due to the nature of INDEL alignment errors: when an INDEL is misaligned, this causes a temporary shift in the read sequence against the reference sequence. This shift manifests as a run of many bases with mismatches due to their incorrect alignment.

We identify regions for realignment by looking for evidence of an insertion or deletion. If a read contains INDEL evidence, we drop a target at the site. We then fold over the targets and merge all targets down into a single target set. The core of our target identification algorithm can be found below.

```
def findTargets(reads: RDD[RichAlignmentRecord],
                maxIndelSize: Int = 500,
                maxTargetSize: Int = 3000): TargetSet = {

  def createTargetSet(target: IndelRealignmentTarget
    ): TargetSet = {
    val tmp = new TreeSet()(TargetOrdering)
    new TargetSet(tmp + target)
  }

  val targets = reads.flatMap(IndelRealignmentTarget(_,
    maxIndelSize))
    .filter(t => !t.isEmpty)

  val targetSet: TargetSet = TargetSet(targets
    .mapPartitions(iter => iter.toArray
      .sorted(TargetOrdering).toIterator)
    .map(createTargetSet)
    .fold(TargetSet())((t1, t2) => joinTargets(t1, t2))
    .set.filter(_.readRange.length <= maxTargetSize))

  targetSet
}
```

We store the targets in Red-Black trees. The performance of the fold depends on the efficiency of the merge function. We achieve efficient merges with the tail-call recursive `join-Targets` function:

```
@tailrec protected final def joinTargets(
  first: TreeSet[IndelRealignmentTarget],
  second: TreeSet[IndelRealignmentTarget]):
    TreeSet[IndelRealignmentTarget] = {

  if (first.isEmpty && second.isEmpty) {
    TreeSet[IndelRealignmentTarget]()(TargetOrdering)
  } else if (second.isEmpty) {
    first
  } else if (first.isEmpty) {
    second
  } else {
    // if the two sets overlap,
    // we must merge their head and tail elements,
    // else we can just blindly append
    if (!TargetOrdering.overlap(first.last, second.head)
      ) {
      first.union(second)
    } else {
      // merge the tail of the first set and the head
      // of the second set and retry the merge
      joinTargets(first - first.last +
        first.last.merge(second.head),
        second - second.head)
    }
  }
}
```

We maintain a sorted-order invariant during the fold. This allows us to know a priori that the elements in the "first" set will always be ordered earlier relative to the elements in the "second" set. However, there can still be overlap between the two sets, as this ordering does not account for the end positions of the targets. If there is overlap between the last target in the "first" set and the first target in the "second" set, we merge these two elements, and try to merge the two sets again. This allows us to find a convex hull of many regions.

The mapping of reads to realignment targets is done through a tail recursive function that performs binary search across the sorted set of realignment targets:

```
@tailrec final def mapToTarget(read: RichAlignmentRecord,
  targets: TreeSet[(IndelRealignmentTarget, Int)]):
    Int = {
  if (targets.size == 1) {
    if (TargetOrdering.contains(targets.head._1, read)) {
      // if there is overlap, return the overlapping target
      targets.head._2
    } else {
      // else, return an empty target (negative index)
      // split up by max indel alignment length for skew
      (-1 - (read.record.getStart / maxTargetLength)).toInt
    }
  } else {
    // split the set and recurse
    val (head, tail) = targets.splitAt(targets.size / 2)
    val reducedSet = if (TargetOrdering.lt(tail.head._1,
        read)) {
      head
    } else {
      tail
    }
    mapToTarget(read, reducedSet)
  }
}
```

This function is applied as a groupBy against all reads. This means that the function is mapped to the RDD that contains all reads. A new RDD is generated where all reads that returned the same INDEL realignment target are grouped together into a list. We randomize reads that do not map to a target across a large number of "fake" targets, which are not processed.

Once we have labeled all reads with a possible target, we generate consensus sequences by inserting INDEL evidence from the reads into the reference genome. For target T, we then pick a consensus sequence C that minimizes the target quality (Q_T):

$$Q_T = \sum_{r \in T} \min(\sum_{i=1}^{L_r} I[r_i = G_i]\epsilon_i, \min_{j \in C} \sum_{i=1}^{L_r} I[r_i = C_{i+j}]\epsilon_i) \quad (2)$$

$I[r_i = x]$ is an indicator function that is 0 if the ith base of the read (r_i) matches the reference or consensus base x. We use G_i to indicate the reference genome base where r_i is original aligned, and C_{i+j} to indicate the $(i + j)$th base of the consensus sequence. The j variable represents the position of the read alignment versus consensus C, and L_r represents the length of read r.

B.3 Duplicate Marking

The following code implements duplicate marking.

```
for (((leftPos, library), readsByLeftPos) <-
    rdd.adamSingleReadBuckets()
      .keyBy(ReferencePositionPair(_))
      .groupBy(leftPositionAndLibrary);

buckets <- {
```

```
leftPos match {
  // These are all unmapped reads.
  // There is no way to determine if
  // they are duplicates
  case None =>
    markReads(readsByLeftPos.unzip._2,
      areDups = false)
  // These reads have their left position mapped
  case Some(leftPosWithOrientation) =>
    // Group the reads by their right position
    val readsByRightPos = readsByLeftPos.groupBy(
      rightPosition)
    // Find any reads with no right position
    val fragments = readsByRightPos.get(None)
    // Check if we have any pairs
    // (reads with a right position)
    val hasPairs = readsByRightPos.keys
      .exists(_.isDefined)
    if (hasPairs) {
      // Since we have pairs,
      // mark all fragments as duplicates
      val processedFrags = if (fragments.isDefined
      ) {
        markReads(fragments.get.unzip._2,
          areDups = true)
      } else {
        Seq.empty
      }
      val processedPairs =
        for (buckets <- (readsByRightPos - None)
          .values;
            processedPair <-
              scoreAndMarkReads(buckets.unzip._2))
        yield processedPair
      processedPairs ++ processedFrags
    } else if (fragments.isDefined) {
      // No pairs. Score the fragments.
      scoreAndMarkReads(fragments.get.unzip._2)
    } else {
      Seq.empty
    }
};
read <- buckets.allReads) yield read
```

First, all reads with the same record group name and read name are collected into buckets. These buckets contain the read and optional mate or secondary alignments. Each read bucket is then keyed by $5'$ position and orientation and grouped together by left (lowest coordinate) position, orientation and library name. Then, we process each read bucket with a common left $5'$ position. Unmapped reads are never marked as duplicate as their position is not known. Mapped reads with a common left position are separated into paired reads and fragments. Fragments, in this context, are reads that have no mate or should have a mate but it doesn't exist.

If there are pairs in a group, all fragments are marked duplicates and the paired reads are grouped by their right $5'$ position. All paired reads that have a common right and left $5'$ position are scored and all but the highest scoring read is marked a duplicate. If there are no pairs in a group, all fragments are scored and all but the highest scoring fragment are marked duplicates.