

# CS286: Database Systems

## 1 Lecture 3—9/4/2014

### 1.1 R\*

- Assumptions:
  - There are administrative causes behind distributed data
  - Network: unreliable transport, in-order, packets are intact
  - Independent node failure
  - Slow-ish network
- Research goals:
  - “Site autonomy”: No centralized state or control
    - \* Data you touch should determine the sites you talk to
    - \* “Distributed system is a system that fails because a machine you’ve never heard of fails”
    - \* Load sharing and decentralization
    - \* Less communication
    - \* Harder to coordinate data consistency
    - \* More network connections beyond hub and spoke
    - \* Metadata management is harder
  - Location transparency → emulate a centralized DB
  - Don’t assume much about the network or OS
- Highlights:
  - Query optimizer cost modeling
  - Data layouts → horizontal partitioning
  - Replication
  - Distribution
  - Query compilation—unclear as to balance between compilation overhead and work saving
  - Spent a lot of time talking about 2PC → presumed commit

### 1.2 Gamma

- Assumptions:
  - Fast interconnect—hypercube, more network bandwidth than aggregate disk bandwidth
  - Shared nothing—no disk or memory sharing
- Research goals:
  - Scale

- Highlights:
  - Parallel hybrid-hash join
  - Chained declustering
- Assess:
  - Linear speedup + scale-up
  - Superlinear speedup due to minimized seek count at scale

## 2 Lecture 4—9/9/2014

- ACID
  - Consistency is not what we typically think
  - Distributed systems: data has a consistent value across sites
  - Databases: data meets contract when transaction completes
- Serializability mathematically gives atomicity and isolation
- Logging gives atomicity and durability
- Ordering:
  - Determines outcome (unless operations are not associative and commutative)
  - Some things are commutable/associable
  - Ordering must be equivalent to some serializable order
  - Implicitly, this provides an API—people don't *need* to reason about concurrency
- What is storage?
  - *Spacial-temporal rendezvous makes everything work!!!!*
- Want to avoid/undo conflicts in space and time
  - *Space*: Shared names
  - *Time*: Ordering
- 2PL: Provides a conflict serialized schedule
  - Ordered by race for locks
  - Ordered by the end of the first phase (“lock point”)
- Multi-version timestamp ordering
  - Every transaction gets a timestamp—this is the only synchronization point
  - For every object:
    - \* Writes generate a new version for an object
    - \* Reads annotate the version for the object

### 3 Lecture 5—9/11/2014

- Good graphs:
  - Crossover points
  - Non-monotonicity
  - Good breadth of X
  - Smooth → variance was accounted for
- Infinite resources:
  - Why run infinite resources? Many people assumed infinite resources in their papers.
  - OCC wins because it allows higher parallelism, at the cost of restarting transactions
  - Blocking (2PL) performs well at start, low at the end. Why?
    - \* Deadlock starts to cause performance to fail
    - \* Lock contention starts to cause transactions to get in each other's way
    - \* Locking is a feedback loop—it lengthens transaction time
- Takeaways:
  - MPL is a control variable—choose your infrastructure for your system
- When do we have “infinite” resources?
  - When we have user interaction (Computer  $\gg$  human)
  - Vastly overprovisioned compute
  - Work is not going on inside the serving infrastructure (e.g., work is done by clients)

#### 3.1 What happens when you go distributed?

- Why go distributed?
  - Capacity (storage and throughput)
  - Low latency (tolerance)
  - Fault tolerance (durability vs. availability)
- Techniques
  - Sharding—split dataset across many nodes
  - Replication

### 4 Lecture 6—9/16/2014

- You need replication → resilience to failure
- Tradeoff between replication and performance
- NoSQL:
  - Typically, a key-value store (data/programming model)
  - Typically distributed and sharded/partitioned
  - Usually weaker consistency model
  - No transactions/weak isolation model
  - “Not MySQL” → lots of work at AOL/etc. with MySQL on memcached

- Typically OSS, not enterprise
- “Scalable”, especially incremental scale → improves organization/administration/ops
- “Evaporation” of the DBA
- Motivations:
  - Bayou: I want to operate when disconnected
  - Dynamo: Nodes gonna fail
- CAP theorem: if partitions occur, then we can either have consistency or availability
  - Availability: As long as a client can access a server, I can access data (concurrent operations don’t need to communicate)
  - Consistency: “linearizable registers” → if I make a write, you can read my write

## 5 Lecture 7—9/18/2014

- In traditional database, have disk page with tuples stored at continuous offsets.
  - Pointers (“slots”) are at end of page and point back to tuples.
  - Can then compress and compact by looking at slot pointers.
  - Fixed length fields stored in tuples
  - Tuples contain pointers to variable length fields
- What changed between 1980 and 2010?
  - CPUs 10,000× faster
  - Disk BW grew 100×
  - Disk seek time improved 10×
- Specifically, gulf between disk performance and processor performance grew
- Research methodology: if area is fairly static, change parameters and see what you can do
- MonetDB
  - Vector/block processing:
- Traditional iterator processing model:
  - Build a tree of operators that run on top of iterators
  - Algorithms have init method (set up state), get next (give me a tuple), and close operators
  - “Pull” model → data and control flow are coupled
- “Late materialization:” query optimizer should defer reading columns until as late as it can
- “Invisible joins:” joins that batch reordering
  - Semijoin: Filter R for all items that have a match in S
- Database cracking: opportunistically reorder blocks in order to improve performance

## 6 Lecture 8—9/23/2014

- Pre-relational data models:
  - Network: objects + pointers
  - Hierarchical: nested sets
- Then, *the Relational Revolution*
- But, persistent questioning:
  - In the 80's, nested relational
  - Object-Oriented Database → 80's/90's
  - And then... XML
- Why flatten into relations:
  - Space efficient
  - Update/delete/inserts require work/care
  - Simple model and language
  - Data independence: physical and logical independence
- When is the relational model a pain?
  - Joins are expensive
  - Must know schema ahead of time
  - Read-only workloads are expensive
  - Programming language state
- Engineering versus “Found Structure” → *bricolage*
  - Engineering → collaboration and communication
  - Found structure → exploratory data analytics
- XML database history:
  - WWW + search ate DB lunch
  - Let's query the internets!
  - $\text{DB} \times \text{WWW} \times \text{markup language people} = \text{pandemonium}$
  - 4 data models, 3 query languages, mostly overlap...
- DB vendors kept up with the pace of research
- How to encode XML:
  - Native
  - Shred to relationable tables
  - Relational encoding of trees
  - Path/value encoding
  - Hybrids

## 7 Lecture 9—9/25/2014

- Few ways to provide/describe isolation; e.g., for a KV store:
  - 2PL → mechanism
  - Avoid anomalies (no lost update, no dirty/fuzzy read) → anomaly prevention
  - Draw conflicts into graph (graph should have no cycle) → graph formalism
  - Every history is view equivalent to a serial history → equivalence formalism
  - If every program preserves an invariant, then every execution will preserve the invariant → integrity
- Full isolation is expensive; let's go for weaker models
- These descriptions don't imply equivalence:
  - Can provide SG without cycle with OCC,  $OCC \neq 2PL$
  - Snapshot isolation → no anomalies, but does not preserve invariants
- Strong vs. weak isolation: weaker implies more anomalies allowed, or more executions allowed
- Conditions of *reality*:
  - Phantom: occurs when your program has a complex predicate, and when table modifications are allowed (modifications can change predicate selections)
    - \* Can solve with predicate locks (but no one does that)
    - \* Can solve with locks on indices, next key locks, etc...
  - *Repeatable read*: in SQL, an isolation level where you can have phantoms, but everything else is OK (IBM defines as full isolation, hence confusion...)
  - *Fuzzy read*: short read locks, I read a single item multiple times and can see different values
  - Statement-level atomicity: hold short read locks for the whole time I'm evaluating a statement
  - `select ... for update`: hint that I'll grab locks later
- Isolation levels:
  0. Short duration write locks
  1. Commit duration write locks
  2. Commit duration write locks, short read locks
  3. Full serializability
- Snapshot isolation: a good demonstration that you can beat anomalies (no lost update, no dirty read, no fuzzy read, no phantoms), but not provide serializability

## 8 Lecture 10—9/30/2014

- Implementation:
  - Concurrency control: 1 transaction at a time
  - All transactions are stored procedures; not interactive
  - Partition your workload (?)
  - Hot replicas: in relational world, replication is a backup strategy, not a runtime strategy
  - Weak consistency “*may*” be interesting

## 8.1 Hekaton

- Built from ground-up to *fit* into SQL Server:
  - Limits design space
  - Must coexist with old systems (e.g., here, must have commit time in old system)
- Much else is in other papers