CS286: Database Systems

1 Lecture 3—9/4/2014

1.1 R*

- Assumptions:
  - There are administrative causes behind distributed data
  - Network: unreliable transport, in-order, packets are intact
  - Independent node failure
  - Slow-ish network
- Research goals:
  - "Site autonomy": No centralized state or control
    - Data you touch should determine the sites you talk to
    - "Distributed system is a system that fails because a machine you've never heard of fails"
    - Load sharing and decentralization
    - Less communication
    - Harder to coordinate data consistency
    - More network connections beyond hub and spoke
    - Metadata management is harder
  - Location transparency → emulate a centralized DB
  - Don't assume much about the network or OS
- Highlights:
  - Query optimizer cost modeling
  - Data layouts → horizontal partitioning
  - Replication
  - Distribution
  - Query compilation—unclear as to balance between compilation overhead and work saving
  - Spent a lot of time talking about 2PC → presumed commit

1.2 Gamma

- Assumptions:
  - Fast interconnect—hypercube, more network bandwidth than aggregate disk bandwidth
  - Shared nothing—no disk or memory sharing
- Research goals:
  - Scale

- Highlights:
  - Parallel hybrid-hash join
  - Chained declustering
- Assess:
  - Linear speedup + scale-up
  - Superlinear speedup due to minimized seek count at scale