CS286: Database Systems

1 Lecture 3—9/4/2014

1.1 R*

- Assumptions:

  - There are administrative causes behind distributed data
  - Network: unreliable transport, in-order, packets are intact
  - Independent node failure
  - Slow-ish network

- Research goals:

  - "Site autonomy": No centralized state or control
    * Data you touch should determine the sites you talk to
    * "Distributed system is a system that fails because a machine you've never heard of fails"
    * Load sharing and decentralization
    * Less communication
    * Harder to coordinate data consistency
    * More network connections beyond hub and spoke
    * Metadata management is harder
  - Location transparency $\rightarrow$ emulate a centralized DB
  - Don't assume much about the network or OS

- Highlights:

  - Query optimizer cost modeling
  - Data layouts $\rightarrow$ horizontal partitioning
  - Replication
  - Distribution
  - Query compilation—unclear as to balance between compilation overhead and work saving
  - Spent a lot of time talking about 2PC $\rightarrow$ presumed commit

1.2 Gamma

- Assumptions:

  - Fast interconnect—hypercube, more network bandwidth than aggregate disk bandwidth
  - Shared nothing—no disk or memory sharing

- Research goals:

  - Scale

- Highlights:
  - Parallel hybrid-hash join
  - Chained declustering
- Assess:
  - Linear speedup + scale-up
  - Superlinear speedup due to minimized seek count at scale

2 Lecture 4—9/9/2014

- ACID
  - Consistency is not what we typically think
  - Distributed systems: data has a consistent value across sites
  - Databases: data meets contract when transaction completes
- Serializability mathematically gives atomicity and isolation
- Logging gives atomicity and durability
- Ordering:
  - Determines outcome (unless operations are not associative and commutative)
  - Some things are commutable/associable
  - Ordering must be equivalent to some serializable order
  - Implicitly, this provides an API—people don't *need* to reason about concurrency
- What is storage?
  - *Spacial-temporal rendezvous makes everything work!!!!*
- Want to avoid/undo conflicts in space and time
  - *Space:* Shared names
  - *Time:* Ordering
- 2PL: Provides a conflict serialized schedule
  - Ordered by race for locks
  - Ordered by the end of the first phase ("lock point")
- Multi-version timestamp ordering
  - Every transaction gets a timestamp—this is the only synchronization point
  - For every object:
    * Writes generate a new version for an object
    * Reads annotate the version for the object

3 Lecture 5—9/11/2014

- Good graphs:
  - Crossover points
  - Non-monotonicity
  - Good breadth of X
  - Smooth $\rightarrow$ variance was accounted for

- Infinite resources:
  - Why run infinite resources? Many people assumed infinite resources in their papers.
  - OCC wins because it allows higher parallelism, at the cost of restarting transactions
  - Blocking (2PL) performs well at start, low at the end. Why?
    * Deadlock starts to cause performance to fail
    * Lock contention starts to cause transactions to get in each other's way
    * Locking is a feedback loop—it lengthens transaction time

- Takeaways:
  - MPL is a control variable—choose your infrastructure for your system

- When do we have "infinite" resources?
  - When we have user interaction (Computer $\gg$ human)
  - Vastly overprovisioned compute
  - Work is not going on inside the serving infrastructure (e.g., work is done by clients)

3.1 What happens when you go distributed?

- Why go distributed?
  - Capacity (storage and throughput)
  - Low latency (tolerance)
  - Fault tolerance (durability vs. availability)

- Techniques
  - Sharding—split dataset across many nodes
  - Replication

4 Lecture 6—9/16/2014

- You need replication $\rightarrow$ resilience to failure

- Tradeoff between replication and performance

- NoSQL:
  - Typically, a key-value store (data/programming model)
  - Typically distributed and sharded/partitioned
  - Usually weaker consistency model
  - No transactions/weak isolation model
  - "Not MySQL" $\rightarrow$ lots of work at AOL/etc. with MySQL on `memcached`

- – Typically OSS, not enterprise
- – "Scalable", especially incremental scale $\rightarrow$ improves organization/administration/ops
- – "Evaporation" of the DBA
- Motivations:
  - – Bayou: I want to operate when disconnected
  - – Dynamo: Nodes gonna fail
- CAP theorem: if partitions occur, then we can either have consistency or availability
  - – Availability: As long as a client can access a server, I can access data (concurrent operations don't need to communicate)
  - – Consistency: "linearizable registers" $\rightarrow$ if I make a write, you can read my write

5 Lecture 7—9/18/2014

- In traditional database, have disk page with tuples stored at continuous offsets.
  - – Pointers ("slots") are at end of page and point back to tuples.
  - – Can then compress and compact by looking at slot pointers.
  - – Fixed length fields stored in tuples
  - – Tuples contain pointers to variable length fields
- What changed between 1980 and 2010?
  - – CPUs 10,000$\times$ faster
  - – Disk BW grew 100$\times$
  - – Disk seek time improved 10$\times$
- Specifically, gulf between disk performance and processor performance grew
- Research methodology: if area is fairly static, change parameters and see what you can do
- MonetDB
  - – Vector/block processing:
- Traditional iterator processing model:
  - – Build a tree of operators that run on top of iterators
  - – Algorithms have init method (set up state), get next (give me a tuple), and close operators
  - – "Pull" model $\rightarrow$ data and control flow are coupled
- "Late materialization:" query optimizer should defer reading columns until as late as it can
- "Invisible joins:" joins that batch reordering
  - – Semijoin: Filter R for all items that have a match in S
- Database cracking: opportunistically reorder blocks in order to improve performance

6   Lecture 8—9/23/2014

- Pre-relational data models:

  - Network: objects + pointers
  - Hierarchical: nested sets

- Then, *the Relational Revolution*

- But, persistent questioning:

  - In the 80's, nested relational
  - Object-Oriented Database $\rightarrow$ 80's/90's
  - And then... XML

- Why flatten into relations:

  - Space efficient
  - Update/delete/inserts require work/care
  - Simple model and language
  - Data independence: physical and logical independence

- When is the relational model a pain?

  - Joins are expensive
  - Must know schema ahead of time
  - Read-only workloads are expensive
  - Programming language state

- Engineering versus "Found Structure" $\rightarrow$ *bricolage*

  - Engineering $\rightarrow$ collaboration and communication
  - Found structure $\rightarrow$ exploratory data analytics

- XML database history:

  - WWW + search ate DB lunch
  - Let's query the internets!
  - DB $\times$ WWW $\times$ markup language people = pandemonium
  - 4 data models, 3 query languages, mostly overlap. . .

- DB vendors kept up with the pace of research

- How to encode XML:

  - Native
  - Shred to relationable tables
  - Relational encoding of trees
  - Path/value encoding
  - Hybrids

7 Lecture 9—9/25/2014

- Few ways to provide/describe isolation; e.g., for a KV store:

  - 2PL → mechanism
  - Avoid anomalies (no lost update, no dirty/fuzzy read) → anomaly prevention
  - Draw conflicts into graph (graph should have no cycle) → graph formalism
  - Every history is view equivalent to a serial history → equivalence formalism
  - If every program preserves an invariant, then every execution will preserve the invariant → integrity

- Full isolation is expensive; let's go for weaker models

- These descriptions don't imply equivalence:

  - Can provide SG without cycle with OCC, OCC $\neq$ 2PL
  - Snapshot isolation → no anomalies, but does not preserve invariants

- Strong vs. weak isolation: weaker implies more anomalies allowed, or more executions allowed

- Conditions of *reality*:

  - Phantom: occurs when your program has a complex predicate, and when table modifications are allowed (modifications can change predicate selections)
    * Can solve with predicate locks (but no one does that)
    * Can solve with locks on indices, next key locks, etc...
  - *Repeatable read*: in SQL, an isolation level where you can have phantoms, but everything else is OK (IBM defines as full isolation, hence confusion...)
  - *Fuzzy read*: short read locks, I read a single item multiple times and can see different values
  - Statement-level atomicity: hold short read locks for the whole time I'm evaluating a statement
  - `select ... for update`: hint that I'll grab locks later

- Isolation levels:

  0. Short duration write locks
  1. Commit duration write locks
  2. Commit duration write locks, short read locks
  3. Full serializability

- Snapshot isolation: a good demonstration that you can beat anomalies (no lost update, no dirty read, no fuzzy read, no phantoms), but not provide serializability

8 Lecture 10—9/30/2014

- Implementation:

  - Concurrency control: 1 transaction at a time
  - All transactions are stored procedures; not interactive
  - Partition your workload (?)
  - Hot replicas: in relational world, replication is a backup strategy, not a runtime strategy
  - Weak consistency "*may*" be interesting

8.1 Hekaton

- Built from ground-up to *fit* into SQL Server:
  - Limits design space
  - Must coexist with old systems (e.g., here, must have commit time in old system)
- Much else is in other papers

9 Lecture 11—10/2/2014

- This guy fellow thinks that people aren't thinking enough about parametrized queries
  - OLTP is pre-canned; OLAP isn't exactly
  - E.g., web forms
- Query optimization, a primer:
  - So, we want to select some keys from a table
  - There are many ways to execute most queries
    * Joins are associative and commutative, so can be reordered
    * Many different algorithms for joins
  - If we have value distribution statistics, we can estimate the cost of a query plan
  - Problem is NP-hard, exponential in the number of tables
  - Getting a plan wrong can cause orders of magnitude performance differences
    * But, there are often many OK plans

9.1 Adaptive Query Processing

- Eddies:
  - Now-sort
  - Led to Inktomi
  - Led to River: adaptive parallel streaming/shuffle
  - Control: approximate query processing
  - Can structure work in many ways:
    * Want to optimize for the expected amount of work

9.2 Robust Query Processing

- How can we make an optimizer more robust?
  - Add a feedback loop and adapt
  - Make query plan "pretty good" even if your estimates are bad—you don't want the *best* plan, you want a good plan that has a very low chance of morphing into an *bad* plan
- Why?
  - Bound your worst-case performance (don't break)
  - If you are robust to a metric, then you can use a cheap approach for collecting that metric
  - Predictable across versions
- Approaches:

- L.C.: You materialized a full intermediate result
- L.C. Eager M.: You add a materialization and look at the full intermediate approach
- E.C.W.C.: Watch a pipe and remember the record IDs, do an anti-join
- E.C.W.B.: Keep a buffer, so not as big as a materialization

- Need dynamic programming algo to search the space

10   Lecture 12—10/7/2014

- Query optimization:
  - We want to find the plan space
  - Estimate the cost!
  - And search (prune) that space
- Should separate a plan space into a logical and physical space
  - Logical space: logical operations
  - Physical space: access operations and disk layout (implementation)
- If you put a rule, you can make things blow up badly
- What happens if an optimizer doesn't exist? You need to take the optimizer that existed before, and modify it for your new optimized implementation.
- Cost estimation:
  - Want to put together summary statistics
  - Specifically, we probably want to generate selectivity estimation for joins
  - Cardinality estimation for group-by's
  - Summary statistics are computed from the data, selectivity/cardinality estimation from summary statistics
  - Cost estimation = statistics
- Logical plan space: SQL, XQuery, Pig, Datalog
- Parallelism:
  - How do I do a parallel database? How do I do joins/etc?
  - How do I allocate resources?
- Physical strategies need to have a logical mapping
- Can define rules mapping logical to logical, physical to logical
- Search strategies:
  - Bottom-up: data oriented, forward chaining
  - Top-down: goal oriented, backward chaining

11 Lecture 13—10/14/2014

- "This is not computer science, it is *sociology*!"

- Ecosystem is important → how does industry impact research?

- **TL; DR:**
  - MapReduce: Introduced mid-query fault tolerance
  - DryadLINQ: "We used cool language shit that we had at Microsoft."

- Early 2000's: What happened?
  - "*The Cloud*" → massive user bases
  - Scale-out data:
    1. Data as a source of value → early offers for FB were based on value of network/data
    2. Data has enormous scale and "schema-on-use" → disk is cheap, don't toss data away!
  - Boom (late '90's) → bust (2001) → boom (2004) → bust (GFC) → now
  - Revenue model changed: money is not made on software sales, rather, it is made on advertising
  - Also, RE: revenue: switched from *sales* to *subscription* → shift from CapEx to OpEx
  - Aside: internet company developers (not enterprise developers) are driving design
  - Custom SW is not sufficiently valuable relative to other things, so give it away:
    1. Network is value, not SW
    2. Ops side arises: even if you *had* FB's software, could you run Facebook?
    3. Open source software attracts top talent
    4. $E[engineer] = 2$ years: if everyone is going to leave in two years, don't have proprietary software
    5. "OOO and software reuse is dead"
  - As a result, *buy vs. build* decision is completely different

- Huge implications on the DB community:
  - *Oracle as a punching bag* → "All the bad things about Oracle must have arrived out of bad DB research."
    * MySQL was also a dichotomous punching bag
    * Where was Postgres? Cool software, but performance is worse than MySQL.
  - We've changed the optimization function: programmer productivity is #1 priority
    * Initially, we must harness parallelism for developers!
      · First, Google MapReduce
      · And, Hadoop @ Yahoo
      · Also, Scope at MSFT
  - Since 2004:
    1. Parallel DBs were overpriced and slow
    2. A lot of the MR workloads were SQL-ish:
       * Hive
       * Pig
    3. Clear OSS opportunity:
       * Cloudera
       * Hortonworks
       * MapR
       * ...
    4. Google sponsors *significant* education push around MapReduce
    5. Hybrid workflows rise up

11.1 Questions

1. Who should use the Hadoop ecosystem?

   - HDFS: Is a pretty useful storage system.
   - Everything else:
     - If you want to program MapReduce (e.g., custom UDFs), MR is a better platform...
     - What do you do once you get a result? With DBs, you can materialize it. With MR, you have many ways to go from there. E.g., *opportunistic vs. agile* pipelines. Assumes sophisticated developers are using the system.
     - How much data should you have? If you've got a program that'll scale across multiple machines, that's a good fit for Hadoop.
     - If you're using Hadoop, your jobs should probably be long running, and not terribly time critical.
     - Machine learning at scale?
     - How does Spark come in? Logical vs. physical fault tolerance is a nice thing to fall back on.

2. What should they pay?

11.2 DryadLINQ and .NET

- LINQ is cool.

- Unfortunately, it is in .NET. . .

- What *is* LINQ? → a functional DSL for collections

- What *in* .NET are they interested in? → Probably libraries for math/collections.

- Debugging: provide a single node implementation that can be run locally.

- Also have PLINQ → parallel LINQ on a single system; so node + processor parallelism

12   Lecture 15—10/21/2014

- Streaming queries was a giant fad:
  - Language semantics issues are a big problem:
    * Need a language that is rich enough that you can use, but
    * Easy enough that you can program
    * Declarative (calculus) vs. algebra languages
    * Complex event processing
  - Whole bucket of systems work:
    * Minimize memory footprint
    * Scaling vs. number/complexity of queries
    * Adaptivity and load shedding
    * *Surprisingly little* work on parallelism and fault tolerance
    * Distributed/service oriented fault tolerance
  - Theory:
    * Using synopses/sketches
    * E.g., what happens if you can only look at a record once?
  - Practice:
    * Streams are not ordered!

          ∗ Therefore, need to *reorder* streams.

          ∗ People like to revoke data.

          ∗ Few streams exist by themselves; we want to join against stored data.

          ∗ Small \$\$\$\$, at least, so far.

- Happened around the same time that XML was a huge fad

- Industry stuff:

  – Wall Street time series stuff (`Kx` for HFT)

  – Enterprise publish/subscribe services

  – Message queues: store and forward networks for transactional endpoints

- Why don't transactions make sense? Concurrent, read heavy workloads.

- Research projects:

  – TelegraphCQ @ Berkeley

  – Aurora/Borealis @ MIT/Brandeis/Brown $\rightarrow$ started as dataflow language

  – STREAM @ Stanford

- Language:

  – Have both data and ordering $\rightarrow$ more than a relation

  – Have many-to-many mapping between records and time

- Operators:

  – Simple monotonic relational operators are fairly easy:

          ∗ Selection

          ∗ Projection

          ∗ Join

  – Nonmonotonic is harder:

          ∗ Temporal windows $\rightarrow$ specifically, when is the window done?

          ∗ Aggregate